

Appendix Real World Data

Laurens van der Maas

12/6/2019

Data Cleaning and Transformation

The code belows cleans and transforms the original dataset from the website.

```
# Large file, code not ran
amsterdam <- read_csv('listingsamsterdam.csv')

amsterdam <- select(amsterdam, price, review_scores_rating, host_since,
                    host_is_superhost, neighbourhood_cleansed, host_listings_count,
                    host_identity_verified, room_type,
                    bathrooms, bedrooms, minimum_nights,
                    number_of_reviews, cancellation_policy, instant_bookable,

# Data transformation
amsterdam$price <- as.numeric(gsub('\\$|', '', amsterdam$price))
amsterdam$weekly_price <- as.numeric(gsub('\\$|', '', amsterdam$weekly_price))
amsterdam$monthly_price <- as.numeric(gsub('\\$|', '', amsterdam$monthly_price))
amsterdam$cleaning_fee <- as.numeric(gsub('\\$|', '', amsterdam$cleaning_fee))

amsterdam$host_since <- as.Date(amsterdam$host_since)
amsterdam <- mutate(amsterdam, host_is_superhost = ifelse(host_is_superhost == TRUE, 1, 0))
amsterdam <- mutate(amsterdam, host_identity_verified = ifelse(host_identity_verified == TRUE, 1, 0))

amsterdam$location_3ways <- ifelse(amsterdam$neighbourhood_cleansed == 'Centrum-West' |
                                   amsterdam$neighbourhood_cleansed == 'Centrum-Oost' |
                                   amsterdam$neighbourhood_cleansed == 'Zuid', "near_centre",
                                   ifelse(amsterdam$neighbourhood_cleansed == 'Bijlmer-Oost' |
                                           amsterdam$neighbourhood_cleansed == 'Gaasperdam - Driemond' |
                                           amsterdam$neighbourhood_cleansed == 'Bijlmer-Centrum' |
                                           amsterdam$neighbourhood_cleansed == 'Osdorp' |
                                           amsterdam$neighbourhood_cleansed == 'Geuzenveld - Slotermee' |
                                           amsterdam$neighbourhood_cleansed == 'Slotervaart' |
                                           amsterdam$neighbourhood_cleansed == 'De Aker - Nieuw Sloten' |
                                           amsterdam$neighbourhood_cleansed == 'Bos en Lommer' |
                                           amsterdam$neighbourhood_cleansed == 'Noord-Oost' |
                                           amsterdam$neighbourhood_cleansed == 'Noord-West' |
                                           amsterdam$neighbourhood_cleansed == 'Oostelijk Havengebied' |
                                           "far_from_centre", "Moderate"))

amsterdam$realprice <- rep(NA, length(amsterdam$price))

for (i in 1:length(amsterdam$realprice)){
  if (amsterdam$minimum_nights[i] > 27){
    if (!is.na(amsterdam$monthly_price[i])){
      amsterdam$realprice[i] <- amsterdam$monthly_price[i]/30
    } else {
      amsterdam$realprice[i] <- amsterdam$price[i]
    }
  }
}
```

```

    }
  } else if (amsterdam$minimum_nights[i] > 6) {
    if (!is.na(amsterdam$weekly_price[i])){
      amsterdam$realprice[i] <- amsterdam$weekly_price[i]/7
    } else {
      amsterdam$realprice[i] <- amsterdam$price[i]
    }
  } else {
    amsterdam$realprice[i] <- amsterdam$price[i]
  }
}

# review_scores_rating has 2565 NA's,
# cleaning_fee 3611 NA's, removing those NAs after removing r_s_r NA's leads to a drop of ~2000 observations
# amsterdam <- filter(amsterdam, !is.na(cleaning_fee))

# 5 NAs in host_since, 5 NAs in host_is_superhost, 5 NAs in host_listingscount
# Host response rate and time have 8536 NA's, removed them
# Experiences_offered contains only none, removed it
# 33 types of property, removed it

# Did not manage to convert host_verifications, removed it
# Minimum nights has maximum of 1001 (outlier I suppose)
# 5 types of cancellation policy
# 7 NAs bathrooms, 14 NAs bedrooms, 8 NAs beds

drops <- c("weekly_price", "monthly_price")
amsterdam <- amsterdam[ , !(names(amsterdam) %in% drops)]

# Michael - Add in host's "age" on AirBnb
Date_scrap <- as.Date("14/09/19", "%d/%m/%y")
amsterdam$host_since_duration <- Date_scrap - amsterdam$host_since

amsterdam <- drop_na(amsterdam)
check <- amsterdam[amsterdam$realprice < 1,] # There is an outlier with price = 0

amsterdam <- amsterdam[amsterdam$realprice > 1,]

amsterdam$logprice <- log(amsterdam$realprice)

```

EDA

```

amsterdam <- read_csv('st443_final_data')

## Warning: Missing column names filled in: 'X1' [1]


## Parsed with column specification:
## cols(
##   X1 = col_double(),
##   review_scores_rating = col_double(),
##   host_is_superhost = col_double(),
##   host_listings_count = col_double(),
##   host_identity_verified = col_double(),

```

```
## room_type = col_character(),
## bathrooms = col_double(),
## bedrooms = col_double(),
## minimum_nights = col_double(),
## number_of_reviews = col_double(),
## cancellation_policy = col_character(),
## instant_bookable = col_logical(),
## cleaning_fee = col_double(),
## location_3ways = col_character(),
## realprice = col_double(),
## host_since_duration = col_double(),
## logprice = col_double()
## )
```

```
# Based on boxplots, there are too many outliers at different price points to be taken out.
# Histograms of numeric variables shows that all are skewed. There is no right way to set limits
# Using logprice transformation however, removes skewness in price. Subsequent OLS shows no outliers up
# Using price as the y variable, plot residuals/leverage and identify 2 outliers,
# they were remove and OLS re-iterated, with more outliers. It will be too long a process.
```


```
# -----
# Price
# Using Laurens ggplots:
ggplot(amsterdam, aes(logprice)) +
  geom_freqpoly(stat='density') + xlim(0,10)
```



```
# log transformation of price result in a less skewed variable "price" without losing any data points
```

```
#boxplot of location vs price
```

```
ggplot(amsterdam, aes(x=location_3ways, y=logprice)) +  
  geom_boxplot(outlier.colour="dark blue", outlier.shape=16,  
              outlier.size=2, notch=TRUE)
```




```
# lesser outliers detected using boxplot of logprice / location
```

```
#boxplot of room_type vs price
```

```
ggplot(amsterdam, aes(x=room_type, y=logprice)) +  
  geom_boxplot(outlier.colour="red", outlier.shape=16, outlier.size=2, notch=TRUE)
```


```
## notch went outside hinges. Try setting notch=FALSE.
```



```
# lesser outliers detected using boxplot of logprice / location

#boxplot of cancellation_policy vs price
ggplot(amsterdam, aes(x=cancellation_policy, y=logprice)) +
  geom_boxplot(outlier.colour="green", outlier.shape=16, outlier.size=2, notch=TRUE)
```


notch went outside hinges. Try setting notch=FALSE.




```
# lesser outliers detected using boxplot of logprice / location

#histograms of numeric variables
par(mfrow=c(2,2))
hist(amsterdam$review_scores_rating)
hist(amsterdam$bathrooms)
hist(amsterdam$bedrooms)
hist(amsterdam$number_of_reviews)
```



histogram of amsterdam\$review_scores_r



Histogram of amsterdam\$bathrooms



Histogram of amsterdam\$bedrooms Histogram of amsterdam\$number_of_reviews



As expected, all are skewed.

#Using log price - interpretation of coefficients will be different

```
logPrice_ols <- lm(logprice ~ review_scores_rating + host_is_superhost +
  host_listings_count + host_identity_verified +
  room_type + bathrooms + bedrooms +
  minimum_nights + number_of_reviews + cancellation_policy +
  instant_bookable + host_since_duration + location_3ways + cleaning_fee, data=amsterdam)
summary(logPrice_ols)
```

##

Call:

```
## lm(formula = logprice ~ review_scores_rating + host_is_superhost +
##   host_listings_count + host_identity_verified + room_type +
##   bathrooms + bedrooms + minimum_nights + number_of_reviews +
##   cancellation_policy + instant_bookable + host_since_duration +
##   location_3ways + cleaning_fee, data = amsterdam)
```

##

Residuals:

```
##      Min       1Q   Median       3Q      Max
## -2.7812 -0.2309 -0.0147  0.2067  4.3489
```

##

Coefficients:

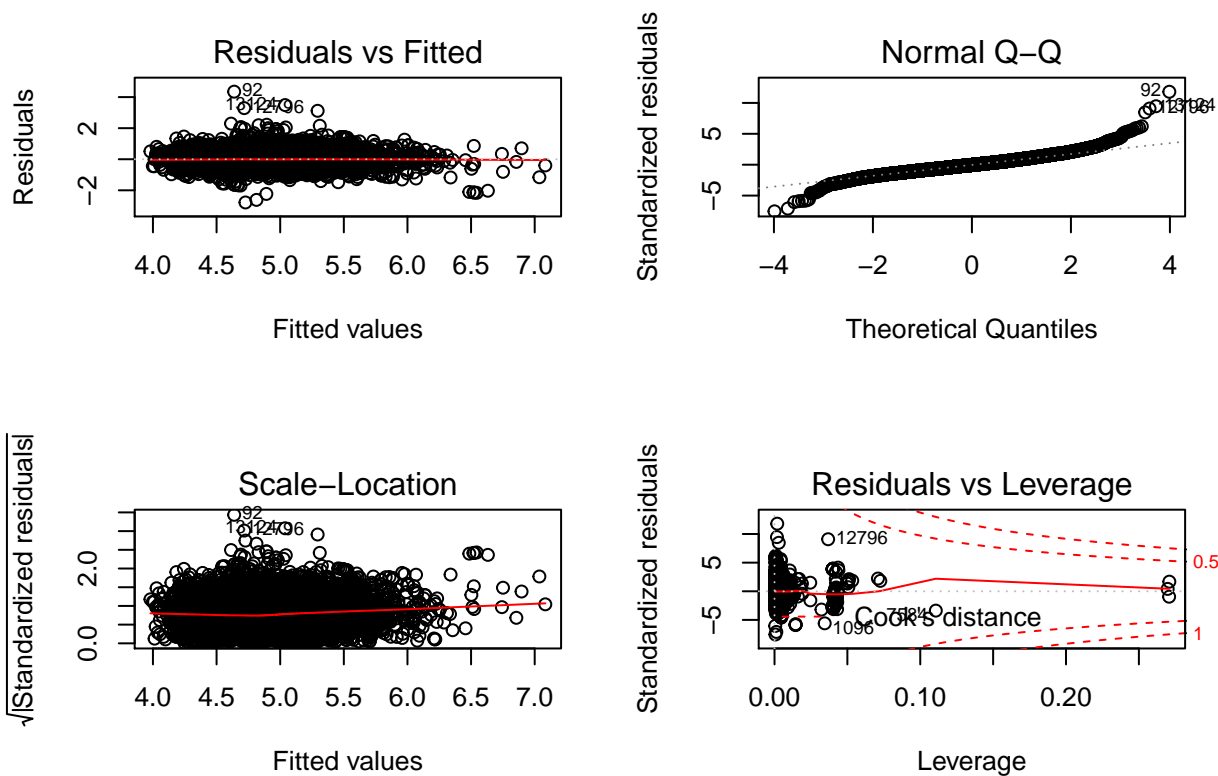
```
##
##              Estimate Std. Error t value
## (Intercept)  3.932e+00  4.868e-02  80.771
## review_scores_rating  3.685e-03  4.915e-04   7.498
## host_is_superhost    9.071e-02  8.623e-03  10.520
## host_listings_count  -5.577e-04  1.070e-04  -5.213
## host_identity_verified -1.146e-03  6.593e-03  -0.174
```

```

## room_typeHotel room          -7.004e-02  2.326e-02  -3.012
## room_typePrivate room        -3.462e-01  9.113e-03 -37.990
## room_typeShared room         -5.522e-01  7.543e-02  -7.320
## bathrooms                    1.079e-01  8.778e-03  12.294
## bedrooms                    1.661e-01  4.019e-03  41.331
## minimum_nights               -2.834e-04  1.922e-04  -1.475
## number_of_reviews            -3.453e-04  6.940e-05  -4.976
## cancellation_policymoderate   1.529e-02  8.966e-03   1.706
## cancellation_policystrict_14_with_grace_period 5.472e-02  8.934e-03   6.125
## cancellation_policysuper_strict_30 6.686e-01  7.370e-02   9.072
## cancellation_policysuper_strict_60 2.900e-01  7.568e-02   3.831
## instant_bookableTRUE         3.068e-02  7.546e-03   4.065
## host_since_duration          -1.888e-05  4.712e-06  -4.006
## location_3waysModerate        1.658e-01  8.051e-03  20.593
## location_3waysnear_centre     3.533e-01  9.163e-03  38.560
## cleaning_fee                 3.468e-03  1.492e-04  23.239
##                               Pr(>|t|)
## (Intercept)                  < 2e-16 ***
## review_scores_rating         6.86e-14 ***
## host_is_superhost            < 2e-16 ***
## host_listings_count          1.88e-07 ***
## host_identity_verified       0.861977
## room_typeHotel room          0.002602 **
## room_typePrivate room        < 2e-16 ***
## room_typeShared room         2.60e-13 ***
## bathrooms                    < 2e-16 ***
## bedrooms                    < 2e-16 ***
## minimum_nights              0.140276
## number_of_reviews            6.58e-07 ***
## cancellation_policymoderate   0.088094 .
## cancellation_policystrict_14_with_grace_period 9.28e-10 ***
## cancellation_policysuper_strict_30 < 2e-16 ***
## cancellation_policysuper_strict_60 0.000128 ***
## instant_bookableTRUE         4.82e-05 ***
## host_since_duration          6.21e-05 ***
## location_3waysModerate        < 2e-16 ***
## location_3waysnear_centre     < 2e-16 ***
## cleaning_fee                 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3688 on 14997 degrees of freedom
## Multiple R-squared:  0.4223, Adjusted R-squared:  0.4216
## F-statistic: 548.2 on 20 and 14997 DF, p-value: < 2.2e-16

```

```
plot(logPrice_ols) # no outliers
```

```
# Descriptive statistics command
lapply(amsterdam, summary)
```

```
## $X1
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1    3755    7510    7510   11264   15018
##
## $review_scores_rating
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  20.00  93.00  97.00   95.09  100.00  100.00
##
## $host_is_superhost
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0000  0.0000  0.0000  0.1794  0.0000  1.0000
##
## $host_listings_count
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.000    1.000    1.000   5.619    1.000  932.000
##
## $host_identity_verified
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##  0.0000  0.0000  0.0000  0.4175  1.0000  1.0000
##
## $room_type
##   Length      Class      Mode
##  15018 character character
##
## $bathrooms
##   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
```

```
## 0.000 1.000 1.000 1.169 1.000 15.000
##
## $bedrooms
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.000 1.000 1.000 1.471 2.000 12.000
##
## $minimum_nights
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 1.000 2.000 2.000 3.294 3.000 1001.000
##
## $number_of_reviews
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.00 5.00 12.00 27.49 27.00 786.00
##
## $cancellation_policy
## Length Class Mode
## 15018 character character
##
## $instant_bookable
## Mode FALSE TRUE
## logical 11342 3676
##
## $cleaning_fee
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 0.00 25.00 39.00 39.07 50.00 500.00
##
## $location_3ways
## Length Class Mode
## 15018 character character
##
## $realprice
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 7.0 100.0 132.0 156.3 180.0 8000.0
##
## $host_since_duration
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 14 1231 1756 1699 2225 4007
##
## $logprice
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## 1.946 4.605 4.883 4.919 5.193 8.987
```

DATA EXPLORATION

```
#amsterdam %>% count(neighbourhood_cleansed) %>% arrange(desc(n)) %>% print(n=30)
amsterdam %>% count(room_type) %>% arrange(desc(n)) %>% print(n=30)
```


```
## # A tibble: 4 x 2
## room_type n
## <chr> <int>
## 1 Entire home/apt 12079
## 2 Private room 2624
## 3 Hotel room 291
## 4 Shared room 24
```

```
amsterdam %>% count(cancellation_policy) %>% arrange(desc(n)) %>% print(n=30)
```


```
## # A tibble: 5 x 2
##   cancellation_policy      n
##   <chr>                <int>
## 1 strict_14_with_grace_period 6621
## 2 moderate                5893
## 3 flexible                2453
## 4 super_strict_30           26
## 5 super_strict_60           25
```

mean review rating is 95... extremely skewed and probably uninteresting density


```
ggplot(amsterdam, aes(review_scores_rating)) +
  geom_freqpoly(stat='density') + xlim(0,100)
```



```
ggplot(amsterdam, aes(bathrooms)) +
  geom_freqpoly(stat='density')
```




```
ggplot(amsterdam, aes(bedrooms)) +  
  geom_freqpoly(stat='density')
```




```
# skewed to the right  
ggplot(amsterdam, aes(number_of_reviews)) +  
  geom_freqpoly(stat='density') + xlim(0,50)
```

```
## Warning: Removed 1853 rows containing non-finite values (stat_density).
```



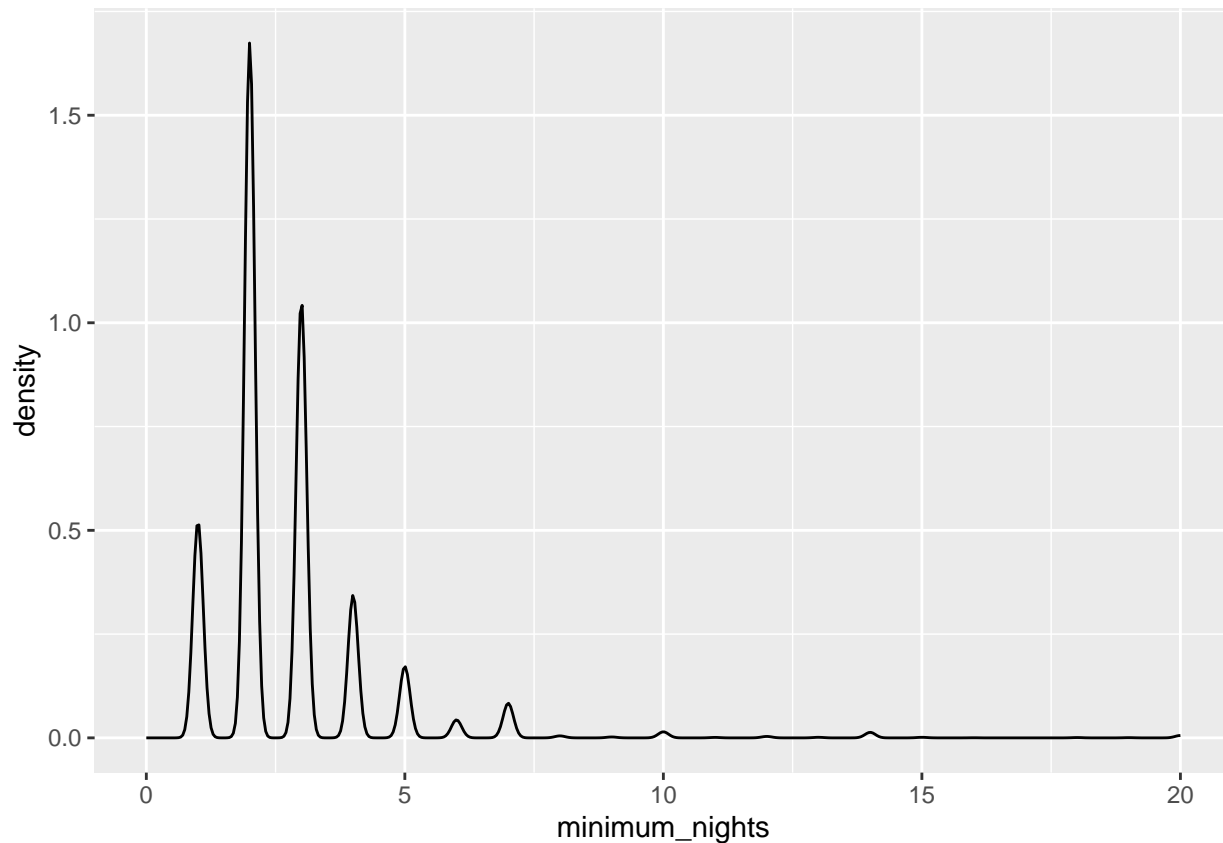
```
# Useless variables, will not add much  
ggplot(amsterdam, aes(host_listings_count)) +  
  geom_freqpoly(stat='density') + xlim(0,20)
```

```
## Warning: Removed 469 rows containing non-finite values (stat_density).
```



```
# Not sure if interesting  
ggplot(amsterdam, aes(minimum_nights)) +  
  geom_freqpoly(stat='density') + xlim(0,20)
```

```
## Warning: Removed 113 rows containing non-finite values (stat_density).
```



```
# make three/four neighbourhoods in terms of price?
#amsterdam %>% group_by(neighbourhood_cleansed) %>%
# summarise(price = mean(price), avgrating = mean(review_scores_rating), n = n()) %>% #arrange(desc(p
```

```
# cancellation relevant interms of price
amsterdam %>% group_by(cancellation_policy) %>%
  summarise(price = mean(logprice), avgrating = mean(review_scores_rating), n = n()) %>% arrange(desc(p
```

```
## # A tibble: 5 x 4
##   cancellation_policy      price avgrating      n
##   <chr>                <dbl>    <dbl> <int>
## 1 super_strict_30        5.66     92.8    26
## 2 super_strict_60        5.51     88.8    25
## 3 strict_14_with_grace_period 4.98     94.9  6621
## 4 moderate              4.89     95.2  5893
## 5 flexible              4.82     95.3  2453
```

```
# room_type significant differences (in terms of price)
amsterdam %>% group_by(room_type) %>%
  summarise(price = mean(logprice), avgrating = mean(review_scores_rating), n = n()) %>% arrange(desc(p
```

```
## # A tibble: 4 x 4
##   room_type      price avgrating      n
##   <chr>        <dbl>    <dbl> <int>
## 1 Entire home/apt 5.01     95.3 12079
## 2 Hotel room     4.94     94.1   291
## 3 Private room   4.52     94.4  2624
## 4 Shared room    4.29     94.5    24
```



```

# instant_bookable provides no information in terms of price and rating
amsterdam %>% group_by(instant_bookable) %>%
  summarise(price = mean(logprice), avgrating = mean(review_scores_rating), n = n()) %>% arrange(desc(p

## # A tibble: 2 x 4
##   instant_bookable price avgrating    n
##   <lgl>           <dbl>     <dbl> <int>
## 1 FALSE           4.93       95.6 11342
## 2 TRUE            4.88       93.6  3676

```

Best Subset Selection, NICE assumptions check

```

#split the data into training and testing dataset
#airbnb1 take out realprice and X columns
#amsterdam <- read_csv('st443_final_data')
airbnb1 = subset(amsterdam, select = -c(1,15))
traingsize = floor(0.7*nrow(airbnb1))
set.seed(123)
train_ind = sample(seq_len(nrow(airbnb1)),size = traingsize)
train=airbnb1[train_ind,]
test=airbnb1[-train_ind,]

airbnb <- amsterdam
attach(airbnb)
str(airbnb)

## Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame': 15018 obs. of  17 variables:
##  $ X1                : num  1 2 3 4 5 6 7 8 9 10 ...
##  $ review_scores_rating : num  98 100 99 95 98 97 80 94 93 97 ...
##  $ host_is_superhost   : num  1 0 1 0 1 0 1 0 0 1 ...
##  $ host_listings_count : num  1 2 1 1 2 1 1 1 1 2 ...
##  $ host_identity_verified: num  0 0 1 1 0 1 0 1 1 0 ...
##  $ room_type           : chr   "Private room" "Entire home/apt" "Private room" "Entire home/apt" ..
##  $ bathrooms           : num  1.5 1 1 1 1 1 1 1.5 1 2 ...
##  $ bedrooms            : num  1 1 1 3 1 2 1 2 1 3 ...
##  $ minimum_nights       : num  3 14 2 3 3 13 30 3 3 6 ...
##  $ number_of_reviews    : num  269 3 200 32 460 690 61 36 3 190 ...
##  $ cancellation_policy  : chr   "strict_14_with_grace_period" "strict_14_with_grace_period" "strict_
##  $ instant_bookable     : logi   TRUE FALSE TRUE FALSE TRUE FALSE ...
##  $ cleaning_fee         : num  60 40 0 60 35 35 50 50 50 0 ...
##  $ location_3ways       : chr   "far_from_centre" "near_centre" "near_centre" "near_centre" ...
##  $ realprice            : num  59 92.9 155 219 159 ...
##  $ host_since_duration  : num  4007 3585 3462 3397 3331 ...
##  $ logprice             : num  4.08 4.53 5.04 5.39 5.07 ...
##  - attr(*, "spec")=
##    .. cols(
##    ..   X1 = col_double(),
##    ..   review_scores_rating = col_double(),
##    ..   host_is_superhost = col_double(),
##    ..   host_listings_count = col_double(),
##    ..   host_identity_verified = col_double(),
##    ..   room_type = col_character(),
##    ..   bathrooms = col_double(),

```

```
## .. bedrooms = col_double(),
## .. minimum_nights = col_double(),
## .. number_of_reviews = col_double(),
## .. cancellation_policy = col_character(),
## .. instant_bookable = col_logical(),
## .. cleaning_fee = col_double(),
## .. location_3ways = col_character(),
## .. realprice = col_double(),
## .. host_since_duration = col_double(),
## .. logprice = col_double()
## .. )
```

```
reg1 = lm(logprice ~., train)
summary(reg1)
```

```
##
## Call:
## lm(formula = logprice ~ ., data = train)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-2.7700	-0.2298	-0.0119	0.2085	3.2918

```
##
## Coefficients:
```


	Estimate	Std. Error	t value
(Intercept)	3.917e+00	5.752e-02	68.099
review_scores_rating	3.741e-03	5.791e-04	6.459
host_is_superhost	8.191e-02	1.027e-02	7.979
host_listings_count	-5.597e-04	1.225e-04	-4.569
host_identity_verified	-2.416e-04	7.783e-03	-0.031
room_typeHotel room	-6.027e-02	2.744e-02	-2.196
room_typePrivate room	-3.397e-01	1.080e-02	-31.446
room_typeShared room	-5.944e-01	8.856e-02	-6.712
bathrooms	1.034e-01	1.106e-02	9.350
bedrooms	1.692e-01	4.893e-03	34.587
minimum_nights	-2.281e-04	1.963e-04	-1.162
number_of_reviews	-3.147e-04	8.172e-05	-3.851
cancellation_policymoderate	1.602e-02	1.061e-02	1.510
cancellation_policystrict_14_with_grace_period	5.426e-02	1.058e-02	5.129
cancellation_policysuper_strict_30	6.286e-01	8.750e-02	7.185
cancellation_policysuper_strict_60	3.237e-01	9.069e-02	3.569
instant_bookableTRUE	3.460e-02	8.959e-03	3.862
cleaning_fee	3.721e-03	1.839e-04	20.236
location_3waysModerate	1.609e-01	9.551e-03	16.847
location_3waysnear_centre	3.456e-01	1.086e-02	31.809
host_since_duration	-1.764e-05	5.621e-06	-3.139

```
## Pr(>|t|)
```

	Pr(> t)
(Intercept)	< 2e-16 ***
review_scores_rating	1.10e-10 ***
host_is_superhost	1.63e-15 ***
host_listings_count	4.96e-06 ***
host_identity_verified	0.975238
room_typeHotel room	0.028087 *
room_typePrivate room	< 2e-16 ***
room_typeShared room	2.01e-11 ***

```
## bathrooms < 2e-16 ***
## bedrooms < 2e-16 ***
## minimum_nights 0.245118
## number_of_reviews 0.000118 ***
## cancellation_policymoderate 0.131051
## cancellation_policystrict_14_with_grace_period 2.97e-07 ***
## cancellation_policysuper_strict_30 7.20e-13 ***
## cancellation_policysuper_strict_60 0.000359 ***
## instant_bookableTRUE 0.000113 ***
## cleaning_fee < 2e-16 ***
## location_3waysModerate < 2e-16 ***
## location_3waysnear_centre < 2e-16 ***
## host_since_duration 0.001699 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3645 on 10491 degrees of freedom
## Multiple R-squared:  0.4259, Adjusted R-squared:  0.4248
## F-statistic: 389.1 on 20 and 10491 DF, p-value: < 2.2e-16
```

```
library(leaps)
reg2=regsubsets(logprice~.,nvmax = 20,data = train)
plot(reg2, scale = "adjr2")
```



```
summary(reg2)
```

```
## Subset selection object
## Call: regsubsets.formula(logprice ~ ., nvmax = 20, data = train)
## 20 Variables (and intercept)
##
## Forced in Forced out
## review_scores_rating FALSE FALSE
## host_is_superhost FALSE FALSE
## host_listings_count FALSE FALSE
## host_identity_verified FALSE FALSE
```

```

## room_typeHotel room                FALSE    FALSE
## room_typePrivate room              FALSE    FALSE
## room_typeShared room               FALSE    FALSE
## bathrooms                         FALSE    FALSE
## bedrooms                         FALSE    FALSE
## minimum_nights                    FALSE    FALSE
## number_of_reviews                 FALSE    FALSE
## cancellation_policymoderate        FALSE    FALSE
## cancellation_policystrict_14_with_grace_period FALSE    FALSE
## cancellation_policysuper_strict_30 FALSE    FALSE
## cancellation_policysuper_strict_60 FALSE    FALSE
## instant_bookableTRUE              FALSE    FALSE
## cleaning_fee                      FALSE    FALSE
## location_3waysModerate             FALSE    FALSE
## location_3waysnear_centre         FALSE    FALSE
## host_since_duration               FALSE    FALSE
## 1 subsets of each size up to 20
## Selection Algorithm: exhaustive
##      review_scores_rating host_is_superhost host_listings_count
## 1  ( 1 ) " "                " "                " "
## 2  ( 1 ) " "                " "                " "
## 3  ( 1 ) " "                " "                " "
## 4  ( 1 ) " "                " "                " "
## 5  ( 1 ) " "                " "                " "
## 6  ( 1 ) " "                " "                " "
## 7  ( 1 ) " "                "*"                " "
## 8  ( 1 ) " "                "*"                " "
## 9  ( 1 ) " "                "*"                " "
## 10 ( 1 ) "*"                "*"                " "
## 11 ( 1 ) "*"                "*"                " "
## 12 ( 1 ) "*"                "*"                " "
## 13 ( 1 ) "*"                "*"                " "
## 14 ( 1 ) "*"                "*"                "*"
## 15 ( 1 ) "*"                "*"                "*"
## 16 ( 1 ) "*"                "*"                "*"
## 17 ( 1 ) "*"                "*"                "*"
## 18 ( 1 ) "*"                "*"                "*"
## 19 ( 1 ) "*"                "*"                "*"
## 20 ( 1 ) "*"                "*"                "*"
##      host_identity_verified room_typeHotel room room_typePrivate room
## 1  ( 1 ) " "                " "                " "
## 2  ( 1 ) " "                " "                "*"
## 3  ( 1 ) " "                " "                "*"
## 4  ( 1 ) " "                " "                "*"
## 5  ( 1 ) " "                " "                "*"
## 6  ( 1 ) " "                " "                "*"
## 7  ( 1 ) " "                " "                "*"
## 8  ( 1 ) " "                " "                "*"
## 9  ( 1 ) " "                " "                "*"
## 10 ( 1 ) " "                " "                "*"
## 11 ( 1 ) " "                " "                "*"
## 12 ( 1 ) " "                " "                "*"
## 13 ( 1 ) " "                " "                "*"
## 14 ( 1 ) " "                " "                "*"

```

```

## 15 ( 1 ) " " " " "*"
## 16 ( 1 ) " " " " "*"
## 17 ( 1 ) " " "*" "*"
## 18 ( 1 ) " " "*" "*"
## 19 ( 1 ) " " "*" "*"
## 20 ( 1 ) "*" "*" "*"
##
## room_typeShared room bathrooms bedrooms minimum_nights
## 1 ( 1 ) " " " " "*" " "
## 2 ( 1 ) " " " " "*" " "
## 3 ( 1 ) " " " " "*" " "
## 4 ( 1 ) " " " " "*" " "
## 5 ( 1 ) " " " " "*" " "
## 6 ( 1 ) " " "*" "*" " "
## 7 ( 1 ) " " "*" "*" " "
## 8 ( 1 ) " " "*" "*" " "
## 9 ( 1 ) "*" "*" "*" " "
## 10 ( 1 ) "*" "*" "*" " "
## 11 ( 1 ) "*" "*" "*" " "
## 12 ( 1 ) "*" "*" "*" " "
## 13 ( 1 ) "*" "*" "*" " "
## 14 ( 1 ) "*" "*" "*" " "
## 15 ( 1 ) "*" "*" "*" " "
## 16 ( 1 ) "*" "*" "*" " "
## 17 ( 1 ) "*" "*" "*" " "
## 18 ( 1 ) "*" "*" "*" " "
## 19 ( 1 ) "*" "*" "*" "*"
## 20 ( 1 ) "*" "*" "*" "*"
##
## number_of_reviews cancellation_policymoderate
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "
## 6 ( 1 ) " " " "
## 7 ( 1 ) " " " "
## 8 ( 1 ) " " " "
## 9 ( 1 ) " " " "
## 10 ( 1 ) " " " "
## 11 ( 1 ) " " " "
## 12 ( 1 ) " " " "
## 13 ( 1 ) "*" " "
## 14 ( 1 ) "*" " "
## 15 ( 1 ) "*" " "
## 16 ( 1 ) "*" " "
## 17 ( 1 ) "*" " "
## 18 ( 1 ) "*" "*"
## 19 ( 1 ) "*" "*"
## 20 ( 1 ) "*" "*"
##
## cancellation_policystrict_14_with_grace_period
## 1 ( 1 ) " "
## 2 ( 1 ) " "
## 3 ( 1 ) " "
## 4 ( 1 ) " "
## 5 ( 1 ) " "

```

```


## 6 ( 1 ) " "
## 7 ( 1 ) " "
## 8 ( 1 ) " "
## 9 ( 1 ) " "
## 10 ( 1 ) " "
## 11 ( 1 ) "*"
## 12 ( 1 ) "*"
## 13 ( 1 ) "*"
## 14 ( 1 ) "*"
## 15 ( 1 ) "*"
## 16 ( 1 ) "*"
## 17 ( 1 ) "*"
## 18 ( 1 ) "*"
## 19 ( 1 ) "*"
## 20 ( 1 ) "*"
##      cancellation_policysuper_strict_30 cancellation_policysuper_strict_60
## 1 ( 1 ) " " " "
## 2 ( 1 ) " " " "
## 3 ( 1 ) " " " "
## 4 ( 1 ) " " " "
## 5 ( 1 ) " " " "
## 6 ( 1 ) " " " "
## 7 ( 1 ) " " " "
## 8 ( 1 ) "*" " "
## 9 ( 1 ) "*" " "
## 10 ( 1 ) "*" " "
## 11 ( 1 ) "*" " "
## 12 ( 1 ) "*" " "
## 13 ( 1 ) "*" " "
## 14 ( 1 ) "*" " "
## 15 ( 1 ) "*" "*"
## 16 ( 1 ) "*" "*"
## 17 ( 1 ) "*" "*"
## 18 ( 1 ) "*" "*"
## 19 ( 1 ) "*" "*"
## 20 ( 1 ) "*" "*"
##      instant_bookableTRUE cleaning_fee location_3waysModerate
## 1 ( 1 ) " " " " " "
## 2 ( 1 ) " " " " " "
## 3 ( 1 ) " " " " " "
## 4 ( 1 ) " " "*" " "
## 5 ( 1 ) " " "*" "*"
## 6 ( 1 ) " " "*" "*"
## 7 ( 1 ) " " "*" "*"
## 8 ( 1 ) " " "*" "*"
## 9 ( 1 ) " " "*" "*"
## 10 ( 1 ) " " "*" "*"
## 11 ( 1 ) " " "*" "*"
## 12 ( 1 ) " " "*" "*"
## 13 ( 1 ) " " "*" "*"
## 14 ( 1 ) "*" "*" "*"
## 15 ( 1 ) "*" "*" "*"
## 16 ( 1 ) "*" "*" "*"
## 17 ( 1 ) "*" "*" "*"

```

```
## 18 ( 1 ) "*"          "*"          "*"
## 19 ( 1 ) "*"          "*"          "*"
## 20 ( 1 ) "*"          "*"          "*"
##      location_3waysnear_centre host_since_duration
## 1 ( 1 ) " "          " "
## 2 ( 1 ) " "          " "
## 3 ( 1 ) "*"          " "
## 4 ( 1 ) "*"          " "
## 5 ( 1 ) "*"          " "
## 6 ( 1 ) "*"          " "
## 7 ( 1 ) "*"          " "
## 8 ( 1 ) "*"          " "
## 9 ( 1 ) "*"          " "
## 10 ( 1 ) "*"         " "
## 11 ( 1 ) "*"         " "
## 12 ( 1 ) "*"         "*"
## 13 ( 1 ) "*"         "*"
## 14 ( 1 ) "*"         " "
## 15 ( 1 ) "*"         " "
## 16 ( 1 ) "*"         "*"
## 17 ( 1 ) "*"         "*"
## 18 ( 1 ) "*"         "*"
## 19 ( 1 ) "*"         "*"
## 20 ( 1 ) "*"         "*"

```

```
plot(reg2, scale = "bic")
```



```
outbs=summary(reg2)
which.max(outbs$adjr2)
```

```
## [1] 19
```

```
which.min(outbs$bic)
```

```
## [1] 16
```

```
#check multicollinearity
library(car)
```

```
## Loading required package: carData
##
## Attaching package: 'car'
## The following object is masked from 'package:dplyr':
##
##      recode
## The following object is masked from 'package:purrr':
##
##      some
```

```
vif(reg1)
```

```
##              GVIF Df GVIF^(1/(2*Df))
## review_scores_rating  1.079635  1      1.039055
## host_is_superhost     1.211551  1      1.100705
## host_listings_count   1.108619  1      1.052910
## host_identity_verified 1.168946  1      1.081178
## room_type             1.445793  3      1.063370
## bathrooms             1.280840  1      1.131742
## bedrooms              1.432958  1      1.197062
## minimum_nights        1.002208  1      1.001104
## number_of_reviews      1.363517  1      1.167697
## cancellation_policy    1.145920  4      1.017172
## instant_bookable       1.173188  1      1.083138
## cleaning_fee           1.340805  1      1.157932
## location_3ways         1.078301  2      1.019025
## host_since_duration    1.210094  1      1.100043
```

```
airbnb2 = subset(train, select = -c(host_identity_verified))
reg3 = lm(logprice~., airbnb2)
vif(reg3)
```

```
##              GVIF Df GVIF^(1/(2*Df))
## review_scores_rating  1.076580  1      1.037584
## host_is_superhost     1.211465  1      1.100666
## host_listings_count   1.101879  1      1.049704
## room_type             1.444603  3      1.063224
## bathrooms             1.280564  1      1.131620
## bedrooms              1.432548  1      1.196891
## minimum_nights        1.002075  1      1.001037
## number_of_reviews      1.357228  1      1.165001
## cancellation_policy    1.144420  4      1.017005
## instant_bookable       1.165138  1      1.079416
## cleaning_fee           1.339441  1      1.157342
## location_3ways         1.078102  2      1.018978
## host_since_duration    1.101082  1      1.049324
```

```
summary(reg3)
```

```
##
## Call:
```



```

## lm(formula = logprice ~ ., data = airbnb2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.7701 -0.2298 -0.0118  0.2084  3.2919
##
## Coefficients:
##                                     Estimate Std. Error t value
## (Intercept)                      3.917e+00  5.746e-02  68.180
## review_scores_rating                3.740e-03  5.783e-04   6.467
## host_is_superhost                   8.191e-02  1.026e-02   7.980
## host_listings_count                -5.594e-04  1.221e-04  -4.581
## room_typeHotel room                -6.026e-02  2.744e-02  -2.196
## room_typePrivate room              -3.397e-01  1.080e-02 -31.452
## room_typeShared room               -5.945e-01  8.855e-02  -6.714
## bathrooms                          1.034e-01  1.106e-02   9.351
## bedrooms                           1.692e-01  4.892e-03  34.593
## minimum_nights                     -2.282e-04  1.962e-04  -1.163
## number_of_reviews                  -3.149e-04  8.153e-05  -3.862
## cancellation_policymoderate         1.602e-02  1.061e-02   1.510
## cancellation_policystrict_14_with_grace_period 5.425e-02  1.058e-02   5.130
## cancellation_policysuper_strict_30        6.286e-01  8.748e-02   7.185
## cancellation_policysuper_strict_60        3.238e-01  9.066e-02   3.571
## instant_bookableTRUE                 3.462e-02  8.927e-03   3.878
## cleaning_fee                       3.721e-03  1.838e-04  20.248
## location_3waysModerate              1.609e-01  9.550e-03  16.848
## location_3waysnear_centre           3.456e-01  1.086e-02  31.810
## host_since_duration                 -1.770e-05  5.361e-06  -3.301
##                                     Pr(>|t|)
## (Intercept)                        < 2e-16 ***
## review_scores_rating               1.04e-10 ***
## host_is_superhost                  1.62e-15 ***
## host_listings_count                4.69e-06 ***
## room_typeHotel room                0.028095 *
## room_typePrivate room              < 2e-16 ***
## room_typeShared room              2.00e-11 ***
## bathrooms                          < 2e-16 ***
## bedrooms                           < 2e-16 ***
## minimum_nights                     0.244919
## number_of_reviews                  0.000113 ***
## cancellation_policymoderate         0.131084
## cancellation_policystrict_14_with_grace_period 2.96e-07 ***
## cancellation_policysuper_strict_30    7.17e-13 ***
## cancellation_policysuper_strict_60    0.000357 ***
## instant_bookableTRUE               0.000106 ***
## cleaning_fee                       < 2e-16 ***
## location_3waysModerate              < 2e-16 ***
## location_3waysnear_centre           < 2e-16 ***
## host_since_duration                0.000967 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3644 on 10492 degrees of freedom
## Multiple R-squared:  0.4259, Adjusted R-squared:  0.4248

```

```
## F-statistic: 409.6 on 19 and 10492 DF, p-value: < 2.2e-16
```

```
#calculate MSE
#predictedvalues = predict(reg3, newdata = test)
#plot(predictedvalues, test$logprice)
#MSE1 = mean((predictedvalues-test$logprice)^2)
##other variable selection method
#step1 = stepAIC(reg1, direction = "both")
#summary(step1)
```

Shrinkage

Ridge

```
## Warning: Missing column names filled in: 'X1' [1]
```

```
## Parsed with column specification:
## cols(
##   X1 = col_double(),
##   review_scores_rating = col_double(),
##   host_is_superhost = col_double(),
##   host_listings_count = col_double(),
##   host_identity_verified = col_double(),
##   room_type = col_character(),
##   bathrooms = col_double(),
##   bedrooms = col_double(),
##   minimum_nights = col_double(),
##   number_of_reviews = col_double(),
##   cancellation_policy = col_character(),
##   instant_bookable = col_logical(),
##   cleaning_fee = col_double(),
##   location_3ways = col_character(),
##   realprice = col_double(),
##   host_since_duration = col_double(),
##   logprice = col_double()
## )
```


```
suppressMessages(library(glmnet))

#amsterdam <- read_csv('st443_final_data')
#amsterdam <- amsterdam[, -c(1,15)]


# glmnet does not use formula language
x <- model.matrix(logprice ~ ., data = amsterdam)
y <- amsterdam$logprice

fit.ridge <- glmnet(x, y, alpha=0)

# 8, 7, 15, 20, 16 most important vars
plot(fit.ridge, xvar="lambda", label= TRUE)
```




```
plot(fit.ridge, xvar="dev", label= TRUE)
```



```
cv.ridge <-cv.glmnet(x, y, alpha=0)
```

```
## Plot of CV mse vs log (lambda), small lambda is best
plot(cv.ridge)
```



```
## Coefficient vector corresponding to the mse which is within
# one standard error of the lowest mse using the best lambda.
coef(cv.ridge)
```

```
## 22 x 1 sparse Matrix of class "dgCMatrix"
##
## (Intercept) 4.0718649364
## (Intercept) .
## review_scores_rating 0.0032651038
## host_is_superhost 0.0652320883
## host_listings_count -0.0003663436
## host_identity_verified -0.0021231936
## room_typeHotel room -0.0222974072
## room_typePrivate room -0.2901882819
## room_typeShared room -0.4627803907
## bathrooms 0.1170837789
## bedrooms 0.1364474682
## minimum_nights -0.0001937531
## number_of_reviews -0.0003570440
## cancellation_policymoderate 0.0035154889
## cancellation_policystrict_14_with_grace_period 0.0434824030
## cancellation_policysuper_strict_30 0.5529793443
## cancellation_policysuper_strict_60 0.2635415827
## instant_bookableTRUE 0.0141006790
## cleaning_fee 0.0035043919
## location_3waysModerate 0.0828561550
## location_3waysnear_centre 0.2367923240
## host_since_duration -0.0000104894
```

```
## Coefficient vector corresponding to the lowest mse using the best lambda
coef(glmnet(x,y,alpha=0, lambda=cv.ridge$lambda.min))
```


```
## 22 x 1 sparse Matrix of class "dgCMatrix"
##
## (Intercept) 3.967646e+00 s0
## (Intercept) .
## review_scores_rating 3.592352e-03
## host_is_superhost 8.430644e-02
## host_listings_count -5.073438e-04
## host_identity_verified -1.472142e-03
## room_typeHotel room -5.567573e-02
## room_typePrivate room -3.328267e-01
## room_typeShared room -5.317271e-01
## bathrooms 1.115002e-01
## bedrooms 1.582633e-01
## minimum_nights -2.600664e-04
## number_of_reviews -3.480520e-04
## cancellation_policymoderate 1.165951e-02
## cancellation_policystrict_14_with_grace_period 5.136956e-02
## cancellation_policysuper_strict_30 6.407773e-01
## cancellation_policysuper_strict_60 2.849320e-01
## instant_bookableTRUE 2.598437e-02
## cleaning_fee 3.518221e-03
## location_3waysModerate 1.400964e-01
## location_3waysnear_centre 3.188280e-01
## host_since_duration -1.652107e-05
```

```
# finding MSE
traingsize = floor(0.7*nrow(amsterdam))
set.seed(123)
train = sample(seq_len(nrow(amsterdam)),size = traingsize)

ridge.train <- glmnet(x[train,], y[train], alpha = 0)
pred.test.ridge <- predict(ridge.train, x[-train,])
dim(pred.test.ridge)
```

```
## [1] 4506 100
```

```
rmse.ridge <- sqrt(apply((y[-train]-pred.test.ridge)^2,2,mean))
plot(log(ridge.train$lambda), rmse.ridge, type="b", xlab="Log(lambda)")
```



```
lambda.best.ridge <- ridge.train$lambda[order(rmse.ridge)[1]]  
lambda.best.ridge
```


```
## [1] 0.02234597
```

```
mseRidge <- min(rmse.ridge)  
mseRidge
```


```
## [1] 0.3795437
```

Lasso

```
fit.lasso <- glmnet(x,y)  
plot(fit.lasso, xvar="lambda", label= TRUE)
```




```
plot(fit.lasso, xvar="dev", label= TRUE)
```



```
cv.lasso <- cv.glmnet(x, y)
# Again, 8, 15, 7, 20.

plot(cv.lasso)
```



```
# Use very small lambda, again
## coefficient vector corresponding to the mse which is within
# one standard error of the lowest mse using the best lambda.
coef(cv.lasso)

## 22 x 1 sparse Matrix of class "dgCMatrix"
##
## (Intercept) 4.165290e+00
## (Intercept) .
## review_scores_rating 2.067228e-03
## host_is_superhost 4.570986e-02
## host_listings_count .
## host_identity_verified .
## room_typeHotel room .
## room_typePrivate room -3.202302e-01
## room_typeShared room -2.559487e-01
## bathrooms 9.280449e-02
## bedrooms 1.586917e-01
## minimum_nights .
## number_of_reviews -5.087404e-05
## cancellation_policymoderate .
## cancellation_policystrict_14_with_grace_period 1.563968e-02
## cancellation_policysuper_strict_30 3.370446e-01
## cancellation_policysuper_strict_60 .
## instant_bookableTRUE .
## cleaning_fee 3.411724e-03
## location_3waysModerate 9.394693e-02
## location_3waysnear_centre 2.738160e-01
## host_since_duration .
```




```
## coefficient vector corresponding to the lowest mse using the best lambda
coef(glmnet(x,y, lambda=cv.lasso$lambda.min))
```

```
## 22 x 1 sparse Matrix of class "dgCMatrix"
##
## (Intercept) 3.940950e+00 s0
## (Intercept) .
## review_scores_rating 3.632109e-03
## host_is_superhost 8.941803e-02
## host_listings_count -5.363794e-04
## host_identity_verified -5.333925e-04
## room_typeHotel room -6.555163e-02
## room_typePrivate room -3.453835e-01
## room_typeShared room -5.437931e-01
## bathrooms 1.074298e-01
## bedrooms 1.658676e-01
## minimum_nights -2.592541e-04
## number_of_reviews -3.362666e-04
## cancellation_policymoderate 1.237059e-02
## cancellation_policystRICT_14_with_grace_period 5.173575e-02
## cancellation_policysuper_strict_30 6.563447e-01
## cancellation_policysuper_strict_60 2.775788e-01
## instant_bookableTRUE 2.921495e-02
## cleaning_fee 3.470366e-03
## location_3waysModerate 1.636307e-01
## location_3waysnear_centre 3.509282e-01
## host_since_duration -1.841013e-05
```

```
## test MSE
lasso.train <-glmnet(x[train,], y[train])
pred.test <-predict(lasso.train, x[-train,])
dim(pred.test)
```

```
## [1] 4506 70
```

```
rmse <-sqrt(apply((y[-train]-pred.test)^2,2,mean))
plot(log(lasso.train$lambda), rmse, type="b", xlab="Log(lambda)")
```



```
lambda.best <- lasso.train$lambda[order(rmse)[1]]
lambda.best
```

```
## [1] 0.0003641836
```

```
mseLasso <- min(rmse)
mseLasso
```

```
## [1] 0.3791684
```

Trees

Codes: Generate training and testing set

```
set.seed(123)
trainingsize <- floor(0.7 * nrow(amsterdam))
trainindex <- sample(seq_len(nrow(amsterdam)), size = trainingsize)
levels(amsterdam$room_type)
```

```
## NULL
```

```
train_df <- amsterdam[trainindex,]
test_df <- amsterdam[-trainindex,]
```

Codes: Decision tree - base model and plots 7 terminal nodes, bedrooms/roomtype+bathroom/location in order of tree hierarchy

```
## Registered S3 method overwritten by 'tree':
##   method      from
##   print.tree cli
## RandomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
##
```


```
## Attaching package: 'randomForest'

## The following object is masked from 'package:dplyr':
##
##      combine

## The following object is masked from 'package:ggplot2':
##
##      margin

## Warning in tree(logprice ~ review_scores_rating + host_is_superhost +
## host_listings_count + : NAs introduced by coercion

##
## Regression tree:
## tree(formula = logprice ~ review_scores_rating + host_is_superhost +
##      host_listings_count + host_identity_verified + room_type +
##      bathrooms + bedrooms + minimum_nights + number_of_reviews +
##      cancellation_policy + instant_bookable + host_since_duration +
##      location_3ways + cleaning_fee, data = train_df)
## Variables actually used in tree construction:
## [1] "bedrooms"      "cleaning_fee"  "bathrooms"
## Number of terminal nodes: 5
## Residual mean deviance: 0.1634 = 1717 / 10510
## Distribution of residuals:
##      Min.    1st Qu.    Median      Mean    3rd Qu.      Max.
## -2.649000 -0.242000 -0.009072  0.000000  0.233600  3.412000
```



Codes: Cross-validation on base decision tree Choose 3 terminal nodes as the decrease in deviation from 3 nodes onwards is minimal.

```
## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by
## coercion


## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion

## Warning in tree(model = m[rand != i, , drop = FALSE]): NAs introduced by
## coercion

## Warning in pred1.tree(tree, tree.matrix(nd)): NAs introduced by coercion
```


Codes: Plot of Prune tree

```
##
## Regression tree:
## snip.tree(tree = tree.final_data_log, nodes = 3L)
## Variables actually used in tree construction:
## [1] "bedrooms"      "cleaning_fee"
## Number of terminal nodes: 3
## Residual mean deviance: 0.1719 = 1807 / 10510
## Distribution of residuals:
##      Min.    1st Qu.    Median      Mean    3rd Qu.     Max.
## -2.649000 -0.242300 -0.009072  0.000000  0.234500  3.412000
```



Codes: Generate predicted value of log price on test_df and calculate MSE (0.1748343)

```
yhat_log <- predict(prune.tree_final_data_log, newdata = test_df)

## Warning in pred1.tree(object, tree.matrix(newdata)): NAs introduced by coercion
tree_final_data_log.test <- test_df[, "logprice"]

## Compute the test MSE
mean((yhat_log - tree_final_data_log.test)^2)

## Warning in mean.default((yhat_log - tree_final_data_log.test)^2): argument is
## not numeric or logical: returning NA

## [1] NA
```

Codes: Use of bagging, m=14 Compare MSE of bagged tree (0.1318838), lower than base decision tree of 0.1748343 Var Imp Plot shows that bedrooms, room type, cleaning fee, locations (in priorities) are the main factors [Note: bedroom is >100%]

```
bag.final_data_log <- randomForest(logprice ~ review_scores_rating + host_is_superhost +
                                     host_listings_count + host_identity_verified +
                                     room_type + bathrooms + bedrooms +
                                     minimum_nights + number_of_reviews + cancellation_policy +
                                     instant_bookable + host_since_duration + location_3ways +
                                     cleaning_fee, data = train_df, mtry=14, importance=TRUE)

bag.final_data_log
```

```
yhat_log.bag <- predict(bag.final_data_log, newdata = test_df)
```

```
## Compute the test MSE
mean((yhat_log.bag - tree_final_data_log.test)^2)
# 0.1318838 MSE
importance(bag.final_data_log)
varImpPlot(bag.final_data_log)
```

Codes: Random forest with $n_{\text{tree}} = 5000$. With 14 features, 3 different random forest models with varying “m” are run ==> $m = \sqrt{14}$, $m = 7$ ($14/2$), and $m = 4$ ($14/3$)

$m = \sqrt{14}$: MSE = 0.1290447 $m = 7$ ($14/2$): MSE = 0.1305521 $m = 4$ ($14/3$): MSE = 0.1290447

```
set.seed(123)
forest.final_data_m1 <- randomForest(logprice ~ review_scores_rating + host_is_superhost +
                                     host_listings_count + host_identity_verified +
                                     room_type + bathrooms + bedrooms +
                                     minimum_nights + number_of_reviews + cancellation_policy +
                                     instant_bookable + host_since_duration + location_3ways +
                                     cleaning_fee, data = train_df, mtry=sqrt(14), importance=TRUE,
                                     n.tree = 5000)

forest.final_data_m1
```

```
## Predicted values on the testing data
yhat.forest_m1 <- predict(forest.final_data_m1, newdata=test_df)
```

```
## Compute the test MSE
mean((yhat.forest_m1 - tree_final_data_log.test)^2)
# MSE of 0.1290447
```

```
set.seed(123)
forest.final_data_m2 <- randomForest(logprice ~ review_scores_rating + host_is_superhost +
                                     host_listings_count + host_identity_verified +
                                     room_type + bathrooms + bedrooms +
                                     minimum_nights + number_of_reviews + cancellation_policy +
                                     instant_bookable + host_since_duration + location_3ways +
                                     cleaning_fee, data = train_df, mtry=7, importance=TRUE,
                                     n.tree = 5000)

forest.final_data_m2
```

```
## Predicted values on the testing data
yhat.forest_m2 <- predict(forest.final_data_m2, newdata=test_df)
```

```
## Compute the test MSE
mean((yhat.forest_m2 - tree_final_data_log.test)^2)
# MSE of 0.1305521
```

```
set.seed(123)
forest.final_data_m3 <- randomForest(logprice ~ review_scores_rating + host_is_superhost +
                                     host_listings_count + host_identity_verified +
                                     room_type + bathrooms + bedrooms +
                                     minimum_nights + number_of_reviews + cancellation_policy +
                                     instant_bookable + host_since_duration + location_3ways +
                                     cleaning_fee, data = train_df, mtry=4, importance=TRUE,
                                     n.tree = 5000)
```

```
forest.final_data_m3

## Predicted values on the testing data
yhat.forest_m3 <- predict(forest.final_data_m3, newdata=test_df)

## Compute the test MSE
mean((yhat.forest_m3 - tree_final_data_log.test)^2)
# MSE of 0.1290447
```

Codes: Boosting with n.tree = 5000. 2 different boosting models with varying depth -> depth=4 and depth=6

Boosting depth = 4: MSE: 0.1383925 ==> relative influence of host_since_duration followed by bedrooms are the highest
 Boosting depth = 6: MSE: 0.1435316 ==> relative influence of host_since_duration followed by bedrooms remains the highest

```
library(gbm)
set.seed(123)
train_df$instant_bookable <- factor(train_df$instant_bookable)
boost.log1 <- gbm( logprice ~ review_scores_rating + host_is_superhost +
  host_listings_count + host_identity_verified +
  room_type + bathrooms + bedrooms +
  minimum_nights + number_of_reviews + cancellation_policy +
  instant_bookable + host_since_duration + location_3ways +
  cleaning_fee, data = train_df, distribution = "gaussian",
  n.trees = 5000, interaction.depth = 4)
summary(boost.log1)

## Predicted values on the testing data
yhat.boost1 <- predict(boost.log1, newdata = test_df, n.trees = 5000)

## Compute the test MSE
mean((yhat.boost1 - tree_final_data_log.test) ^ 2)
#MSE of 0.1383925
```

```
set.seed(123)
train_df$instant_bookable <- factor(train_df$instant_bookable)
boost.log2 <- gbm( logprice ~ review_scores_rating + host_is_superhost +
  host_listings_count + host_identity_verified +
  room_type + bathrooms + bedrooms +
  minimum_nights + number_of_reviews + cancellation_policy +
  instant_bookable + host_since_duration + location_3ways +
  cleaning_fee, data = train_df, distribution = "gaussian",
  n.trees = 5000, interaction.depth = 6)
summary(boost.log2)

## Predicted values on the testing data
yhat.boost2 <- predict(boost.log2, newdata = test_df, n.trees = 5000)

## Compute the test MSE
mean((yhat.boost2 - tree_final_data_log.test) ^ 2)
#MSE of 0.1435316
```

```
set.seed(123)
train_df$instant_bookable <- factor(train_df$instant_bookable)
```

```

boost.log3 <- gbm( logprice ~ review_scores_rating + host_is_superhost +
                  host_listings_count + host_identity_verified +
                  room_type + bathrooms + bedrooms +
                  minimum_nights + number_of_reviews + cancellation_policy +
                  instant_bookable + host_since_duration + location_3ways +
                  cleaning_fee, data = train_df, distribution = "gaussian",
                  n.trees = 5000, interaction.depth = 2)
summary(boost.log3)

## Predicted values on the testing data
yhat.boost3 <- predict(boost.log3, newdata = test_df, n.trees = 5000)

## Compute the test MSE
mean((yhat.boost3 - tree_final_data_log.test) ^ 2)
#MSE of 0.1318706

```

GAM

Code below.

```

library("gam")

poly1 = lm(logprice~poly(bedrooms,4), data = airbnb1)
summary(poly1)
poly2 = lm(logprice~poly(bathrooms,3),data = airbnb1)
summary(poly2)
poly3 = lm(logprice ~ poly(number_of_reviews,4),data = airbnb1)
summary(poly3)
plot(bathrooms, logprice)

#gam1 is trying natural spline
gam1 = lm(logprice ~ ns(bedrooms,4)+ns(bathrooms,2)+review_scores_rating+host_is_superhost+host_listings)
summary(gam1)
#gam1p is trying smooth spline
gam1p = lm(logprice ~ s(bedrooms,4)+s(bathrooms,2)+review_scores_rating+host_is_superhost+host_listings)
summary(gam1p)

bestgam = regsubsets(logprice ~ ns(bedrooms,4)+ns(bathrooms,2)+review_scores_rating+host_is_superhost+host_listings)
plot(bestgam, scale = "adjr2")

gam2 = lm(logprice ~ ns(bedrooms,4)+ns(bathrooms,2)+review_scores_rating+host_is_superhost+host_listings)
summary(gam2)
#calculate MSE(gam)
predictedvalues1 = predict(gam2, newdata = test)
plot(predictedvalues1, test$logprice)
MSE2 = mean((predictedvalues1-test$logprice)^2)
#We can see that the MSE is around 0.1338, slightly worse than the previous method.

gam3 = lm(logprice ~ ns(bedrooms,4), test)
plot(gam3)
shapiro.test(gam3$residuals)

```


Neural Networks

```
# NN - Part 1: Data transformation
#amsterdam <- read.csv("st445_final_data", header = T)

amsterdam <- amsterdam[,-1]
#dummify the data
amsterdam <- mutate(amsterdam,
                     instant_bookable = ifelse(instant_bookable == TRUE, 1, 0))
#output_vector = amsterdam[, 'logprice']
amsterdam <- fastDummies::dummy_cols(amsterdam)
#amsterdam <- amsterdam[,-c(5,10,13,14,22,26)]
amsterdam <- amsterdam[,-c(5,10,13,14)]

# Set training and testing dataset
set.seed(123)
trainingsize <- floor(0.7 * nrow(amsterdam))
trainindex <- sample(seq_len(nrow(amsterdam)), size = trainingsize)

train_df <- amsterdam[trainindex,]
test_df <- amsterdam[-trainindex,]

# split up train features(x) and train targets(y)
train_data <- as.matrix(train_df[,-12])
train_targets <- as.array(train_df[,12])

# split up test features(x) and test targets(y)
test_data <- as.matrix(test_df[,-12])
test_targets <- as.array(test_df[,12])

#Scale the data so that all variables are between 0 and 1

mean <- apply(train_data, 2, mean)
std <- apply(train_data, 2, sd)
train_data <- scale(train_data, center = mean, scale = std)
test_data <- scale(test_data, center = mean, scale = std)

# NN - PArt 2: Build neural network model
build_model <- function() {
  model <- keras_model_sequential() %>%
    layer_dense(units = 16, activation = "relu",
                 input_shape = dim(train_data)[[2]]) %>%
    layer_dense(units = 16, activation = "relu") %>%
    layer_dense(units = 1) # single node because it is a regression ML

  model %>% compile(
    optimizer = "rmsprop",
    loss = "mse",
    metrics = c("mae")
  )
}

# NN - Part 3 - change the number of learning iterations, i.e "num_epochs", with 10 and 50
all_scores <- c()
```

```

num_epochs <- 50

# Build the Keras model (already compiled)
model <- build_model()
summary(model)
# Train the model (in silent mode, verbose=0)
model %>% fit(train_data, train_targets,
             epochs = num_epochs, batch_size = 1, verbose = 0)

# Evaluate the model on the validation data
results <- model %>% evaluate(test_data, test_targets, verbose = 0)
#all_scores <- c(all_scores, results$mean_absolute_error)
all_scores <- c(all_scores, results$loss)
#MSE: 0.149141778 using num_epoch = 10, batch size=1
#MSE: 0.15116772 using num_epoch = 50, batch size=1

```

XGBoost

```

library(dplyr)
library(xgboost)

##
## Attaching package: 'xgboost'

## The following object is masked from 'package:dplyr':
##
##     slice

library(stringr)
library(caret)

## Loading required package: lattice

##
## Attaching package: 'caret'

## The following object is masked from 'package:purrr':
##
##     lift

library(car)
library(fastDummies)
library(ModelMetrics)

##
## Attaching package: 'ModelMetrics'

## The following objects are masked from 'package:caret':
##
##     confusionMatrix, precision, recall, sensitivity, specificity

## The following object is masked from 'package:base':
##
##     kappa

```

```

#amsterdam <- read_csv('st443_final_data')
#amsterdam <- amsterdam[, -c(1,15)]

amsterdam <- mutate(amsterdam,
                    instant_bookable = ifelse(instant_bookable == TRUE, 1, 0))

amsterdam <- fastDummies::dummy_cols(amsterdam)
amsterdam <- amsterdam[, -c(5,10,13,16,22,26)]

traingsize = floor(0.7*nrow(amsterdam))
set.seed(123)
trainindex = sample(seq_len(nrow(amsterdam)), size = traingsize)

train_df <- amsterdam[trainindex,]
test_df <- amsterdam[-trainindex,]

trainmatrix <- as.matrix(train_df, rownames.force = NA)
testmatrix <- as.matrix(test_df, rownames.force = NA)
dtrain <- as(trainmatrix, "sparseMatrix")
dtest <- as(testmatrix, "sparseMatrix")

train_data <- xgb.DMatrix(data = dtrain[, -12], label = dtrain[, "logprice"])
test_data <- xgb.DMatrix(data = dtest[, -12])

xgb_grid = expand.grid(
  nrounds = 1000,
  eta = c(0.1, 0.05, 0.01),
  max_depth = c(2, 3, 4, 5, 6),
  gamma = 0,
  colsample_bytree=1,
  min_child_weight=c(1, 2, 3, 4, 5),
  subsample=1
)

my_control <- trainControl(method="cv", number=5)

# Not run, takes ages
#xgb_caret <- train(x = train_df[-12], y = train_df$logprice,
#                  method='xgbTree', trControl= my_control,
#                  tuneGrid = xgb_grid)
#xgb_caret$bestTune
# nrounds = 1000, max_depth = 5, eta = 0.01, min_child_weight = 1

#xgb_tune <- train(logprice ~.,
#                 data = train_df,
#                 method="xgbLinear",
#                 metric = "RMSE",
#                 trControl = cv.ctrl,
#                 tuneGrid = xgb_grid
#)

default_param <- list(

```

```

objective = "reg:linear",
booster = "gbtree",
eta=0.01, #default = 0.3
gamma=0,
max_depth=5, #default=6
min_child_weight=1, #default=1
subsample=1,
colsample_bytree=1
)

#xgbcv <- xgb.cv( params = default_param,
#                data = dtrain, nrounds = 2000,
#                nfold = 5,
#                showsd = T,
#                stratified = T,
#                print_every_n = 40,
#                early_stopping_rounds = 10,
#                maximize = F,
#                label = dtrain[, "logprice"])

xgb_mod <- xgb.train(data = train_data, params = default_param, nrounds = 1300)
XGBpred <- predict(xgb_mod, test_data)
rmse <- rmse(test_df$logprice, XGBpred)

library(Ckmeans.1d.dp) #required for ggplot clustering
mat <- xgb.importance(feature_names = colnames(train_df[-12]), model = xgb_mod)
xgb.ggplot.importance(importance_matrix = mat[1:20], rel_to_first = TRUE)

```

